

Probabilistic Artificial Intelligence


Introduction to Reinforcement Learning

Prof. Andreas Krause
Learning and Adaptive Systems (las.ethz.ch)

Announcements

- No lecture next week (13.12.), only exercises

Learning to Act in Unknown Environments



Action	Reward
Forward	0
Left	0
Forward	0
Right	0
...	
Forward	10

- Learn a mapping from (seq. of) actions to rewards
- **Credit assignment problem:** which actions got me to the large reward?

Reinforcement learning

Agent actions *change* the state of the world (in contrast to supervised learning)

World: “You are in state x_{17} . You can take actions a_3 and a_9 ”

Agent: “I take a_3 ”

World: “You get reward -4 and are now in state x_{279} . You can take actions a_7 and a_9 ”


Agent: “I take a_9 ”

World: “You get reward 27 and are now in state $x_{279} \dots$ You can take actions a_2 and a_{17} ”

...


Assumption: States change according to some (unknown) MDP!

RL = Planning in unknown MDPs



Solving the Credit Assignment Problem

State x	Action a	Reward r
PU	A	0
PU	S	0
PU	A	0
PF	S	0
PF	A	10
PF	A	10
...




Observed state transitions and rewards let you **learn** the underlying MDP!

Recall: Planning in MDPs

- Deterministic policy $\pi: X \rightarrow A$
- Induces a **Markov chain**: $X_0, X_1, \dots, X_t, \dots$ with transition probabilities

$$P(X_{t+1}=x' | X_t=x) = P(x' | x, \pi(x))$$



- Expected value $J(\pi) = E[r(X_0, \pi(X_0)) + \gamma r(X_1, \pi(X_1)) + \gamma^2 r(X_2, \pi(X_2)) + \dots]$
- **Value function**:

$$V^\pi(x) = J(\pi | X_0 = x) = \mathbb{E} \left[\sum_{t=0}^{\infty} \gamma^t r(X_t, \pi(X_t)) | X_0 = x \right]$$

Value Functions and Policies

Every value function induces a policy

Value function V^π

$$V^\pi(x) = r(x, \pi(x)) + \gamma \sum_{x'} P(x' | x, \pi(x)) V^\pi(x')$$

Greedy policy w.r.t. V

$$\pi_V(x) = \operatorname{argmax}_a r(x, a) + \gamma \sum_{x'} P(x' | x, a) V(x')$$

Every policy induces a value function

Theorem (Bellman):

Policy optimal \Leftrightarrow greedy w.r.t. its induced value function!

$$V^*(x) = \max_a [r(x, a) + \gamma \sum_{x'} P(x' | x, a) V^*(x')]$$

Reinforcement Learning

- RL is different from supervised learning
 - The data we get is not i.i.d.
 - In reinforcement learning, the data we get *depends on our actions!*
 - Some actions have higher rewards than others!
- *Exploration—Exploitation Dilemma*: Should we
 - **Explore**: gather more data to avoid missing out on a potentially large reward?
 - **Exploit**: stick with our current knowledge and build an optimal policy for the data we've seen?

Two basic approaches to RL

1) Model-based RL


- Learn the MDP
 - Estimate transition probabilities $P(x' | x, a)$
 - Estimate reward function $r(x, a)$
- Optimize policy based on estimated MDP

2) Model-free RL

- Estimate the value function directly;
- Policy gradient methods;
- Actor-critic methods.

Off-policy vs on-policy RL

- On-policy RL
 - Agent has full control over which actions to pick
 - Can choose how to trade exploration and exploitation
- Off-policy RL
 - Agent has no control over actions, only gets observational data (e.g., demonstrations, data collected by applying a different policy, ...)



Learning the MDP

- Need to estimate
 - transition probabilities $P(X_{t+1} | X_t, A)$
 - Reward function $r(X, A)$
- Can use techniques from learning Bayes Nets:
(regularized) maximum likelihood estimation

- Data set: $X_1, a_1, r_1, X_2, a_2, r_2, X_3, a_3, r_3, X_4, \dots$
 $\hookrightarrow D = \{(x_1, a_1, r_1, x_2), (x_2, a_2, r_2, x_3), (x_3, a_3, r_3, x_4), \dots\}$

- Estimate transitions:

$$P(X_{t+1} | X_t, A) \approx \frac{\text{Count}(X_{t+1}, X_t, A)}{\text{Count}(X_t, A)}$$

- Estimate rewards:

$$r(x, a) \approx \frac{1}{N_{x,a}} \sum_{t: X_t=x, A_t=a} R_t$$

↑ rewards experienced

Exploration-Exploitation Dilemma

- Always pick a random action?
 - Will eventually correctly estimate all probabilities and rewards 😊
 - May do extremely poorly in terms of rewards! 😞
- Always pick the best action according to current knowledge?
 - Quickly get some reward
 - Can get stuck in suboptimal action! 😞
- Balance exploration and exploitation (more later)

Trading Exploration and Exploitation

- ε_t greedy
 - With probability ε_t : Pick random action
 - With probability $(1-\varepsilon_t)$: Pick best action
- If sequence ε_t satisfies Robbins Monro (RM) conditions then will converge to optimal policy with probability 1

$$\sum_t \varepsilon_t = \infty ; \quad \sum_t \varepsilon_t^2 < \infty$$

- Simple, often performs fairly well
- Doesn't quickly eliminate clearly suboptimal actions

The R_{\max} Algorithm [Brafman & Tennenholtz '02]

Optimism in the face of uncertainty!


- If you don't know $r(x, a)$:
 - Set it to R_{\max} !
- If you don't know $P(x' | x, a)$:
 - Set $P(x^* | x, a) = 1$ where x^* is a “fairy tale” state:

$$r(x^*, a) = R_{\max}$$
$$P(x^* | x, a) = 1 \quad \forall a$$

stay at that state

Implicit Exploration Exploitation in R_{\max}

1st 2nd 3rd



Three actions:

- Left
- Right
- Dig

$$r(i, \text{Left}) = 0$$

$$r(i, \text{Right}) = 0$$

Never need to explicitly choose whether we're exploring or exploiting!

Can rule out clearly suboptimal actions very quickly

Exploration—Exploitation Lemma

with high probability

Theorem: Every T timesteps, w.h.p., R_{\max} either

- Obtains near-optimal reward, or
- Visits at least one unknown state-action pair

- T is related to the mixing time of the Markov chain of the MDP induced by the optimal policy

The R_{\max} algorithm

Input: Starting state x_0 , discount factor γ

Initially:

- Add fairy tale state x^* to MDP
- Set $r(x, a) = R_{\max}$ for all states x and actions a
- Set $P(x^* | x, a) = 1$ for all states x and actions a
- Choose optimal policy for r and P

Repeat:

- Execute policy π
- For each visited state action pair x, a , update $r(x, a)$
- Estimate transition probabilities $P(x' | x, a)$
- If observed “enough” transitions / rewards, recompute policy π according to current model P and r

How much is “enough”?

How many samples do we need to accurately estimate
 $P(x' | x, a)$ or $r(x, a)$?

Hoeffding-Chernoff bound:

- X_1, \dots, X_n i.i.d. samples from Bernoulli distribution w. mean μ

$$P\left(\left|\mu - \frac{1}{n} \sum_i X_i\right| \geq \varepsilon\right) \leq 2 \exp(-2n\varepsilon^2)$$

Performance of R_{\max} [Brafman & Tennenholz]

Theorem:

With probability $1-\delta$, R_{\max} will reach an ε -optimal policy in a number of steps that is polynomial in $|X|$, $|A|$, T , $1/\varepsilon$ and $\log(1/\delta)$

Problems of model-based RL?

- Memory required:
store $P(x'|x,a)$ $\Rightarrow O(|X|^2 \cdot |A|)$
 $r(x,a)$ $\Rightarrow O(|X| \cdot |A|)$
- Computation time:

Solving MDP once: $\text{poly}(|X|, |A|, \frac{1}{\varepsilon}, \log \delta')$
Need to do this often!

Two basic approaches

1) Model-based RL

- Learn the MDP
 - Estimate transition probabilities $P(x' | x, a)$
 - Estimate reward function $r(x, a)$
- Optimize policy based on estimated MDP

2) Model-free RL

- Estimate the value function directly;
- Policy gradient methods;
- Actor-critic methods

Model free RL

- Recall:
 1. Optimal value function $V^*(x) \rightarrow$ opt. policy π^*
 2. For optimal value function it holds:

$$V^*(x) = \max_a Q^*(x, a)$$

$$\text{where } \underline{Q^*(x, a)} = r(x, a) + \gamma \sum_{x'} P(x' \mid x, a) V^*(x')$$

Key idea: Estimate $Q^*(x, a)$ directly from samples!

Q-learning

- Estimate $Q^*(x, a) = r(x, a) + \gamma \sum_{x'} P(x' | x, a) V^*(x')$

$$V^*(x) = \max_a Q^*(x, a)$$

$$\pi^*(x) = \arg\max_a Q^*(x, a)$$

- Suppose we
 - Have initial estimate of $Q(x, a)$
 - observe transition x, a, x' with reward r

$$Q(x, a) \leftarrow (1 - \alpha_t)Q(x, a) + \alpha_t \left(r + \gamma \max_{a'} Q(x', a') \right)$$

$V(x')$

Q-learning

$$Q(x, a) \leftarrow (1 - \alpha_t)Q(x, a) + \alpha_t \left(r + \gamma \max_{a'} Q(x', a') \right)$$

Theorem: If learning rate α_t satisfies

$$\sum_t \alpha_t = \infty$$

$$\sum_t \alpha_t^2 < \infty$$

and actions are chosen at random, then Q learning converges to optimal Q^* with probability 1

How can we trade off exploration and exploitation?

Convergence of Optimistic Q-learning

[Even-dar & Mansour '02]

Similar to R_{\max} :

Initialize $Q(x, a) = \frac{R_{\max}}{1 - \gamma} \prod_{t=1}^{T_{\text{init}}} (1 - \alpha_t)^{-1}$

Theorem: With prob. $1 - \delta$, optimistic Q -learning obtains an ε -optimal policy after a number of time steps that is polynomial in $|X|$, $|A|$, $1/\varepsilon$ and $\log(1/\delta)$

Properties of Q-learning

- Memory required: store $Q(x, a)$ $\Rightarrow O(|x||A|)$
- Computation time:
At each iter.,
pick $a_f = \underset{a}{\operatorname{argmax}} Q(x, a) \Rightarrow O(|A|)$
update $Q(x, a)$ $\Rightarrow O(|A|)$

Challenges of RL

- MDP and RL polynomial in $|A|$ and $|X|$. Problem in:
 - Structured domains (chess, multiagent planning, ...):
 $|X|, |A|$ exponential in #agents, state variables, ...
 - Continuous domains ($|A|$ and $|X|$ infinite)
 - ➔ Learning / approximating value functions (regression)
 - ➔ Approximate planning using factored representations
-
- Risk in exploration
 - Random exploration can be disastrous
 - ➔ Learn from “safe” examples: Apprenticeship learning

Acknowledgments

- Slides based on material accompanying the textbook “AI: A Modern Approach” (3rd edition) by S. Russell and P. Norvig and material by Carlos Guestrin