

# Probabilistic AI Tutorial

## Bayesian Learning and RL

---

Vincent Fortuin (*fortuin@inf.ethz.ch*)

December 2019

Institute of Machine Learning, ETH Zürich

## TABLE OF CONTENTS

### 1. Bayesian learning

Learning as inference

Bayesian linear regression

Gaussian processes

Bayesian deep learning

### 2. Reinforcement learning

Types of RL

Q-learning

## Bayesian learning

---

- Let us have a data set  $\mathbf{X} = [\mathbf{x}_1, \dots, \mathbf{x}_n]^\top \in \mathbb{R}^{n \times d}$  with labels  $\mathbf{y} = [y_1, \dots, y_n]^\top \in \mathbb{R}^n$ .
- We assume there is an underlying function  $f : \mathbb{R}^d \rightarrow \mathbb{R}$  with  $y = f(\mathbf{x}) + \epsilon$  where  $\epsilon$  is some noise.
- We are trying to approximate  $f$  with some parameterized function  $\hat{f}_\theta$ , such that  $\hat{f}_\theta(\mathbf{x}) \approx f(\mathbf{x})$ .

## OPTIMIZATION VS. INFERENCE

Optimization view:

$$\theta^* = \arg \min_{\theta} \mathcal{L}(\theta; \mathbf{X}, \mathbf{y}) \quad \text{with} \quad \mathcal{L}(\theta; \mathbf{X}, \mathbf{y}) = \sum_{i=1}^n D(\hat{f}_{\theta}(\mathbf{x}_i), y_i) + \lambda \Omega(\theta)$$

for discrepancy  $D(\cdot, \cdot)$  and regularizer  $\Omega(\cdot)$

Inference view:

$$p(\theta^* \mid \mathbf{X}, \mathbf{y}) = \frac{p(\mathbf{X}, \mathbf{y} \mid \hat{f}_{\theta^*}) p(\theta^*)}{p(\mathbf{X}, \mathbf{y})} = \frac{p(\mathbf{y} \mid \mathbf{X}, \hat{f}_{\theta^*}) p(\theta^*)}{\sum_{\tilde{\theta}} p(\mathbf{y} \mid \mathbf{X}, \hat{f}_{\tilde{\theta}})}$$

- Assume  $y = f(\mathbf{x}) + \epsilon$  where  $\epsilon \sim \mathcal{N}(0, \sigma_y^2)$ .
- The log likelihood is

$$\log p(\mathbf{y} | \mathbf{X}, \hat{f}_\theta) = \log \mathcal{N}(\mathbf{y}; \hat{f}_\theta(\mathbf{X}), \sigma_y^2 \mathbf{I}) = -\frac{1}{2\sigma_y^2} \|\hat{f}_\theta(\mathbf{X}) - \mathbf{y}\|_2^2 + \text{const.}$$

- With a linear model  $\hat{f}_\theta(\mathbf{x}) = \mathbf{x}^\top \theta$  this gives

$$\theta_{MLE} = \arg \max_{\theta} p(\mathbf{y} | \mathbf{X}, \hat{f}_\theta) = \arg \min_{\theta} \|\hat{f}_\theta(\mathbf{X}) - \mathbf{y}\|_2^2$$

- This can be solved in closed form<sup>1</sup> as

$$\theta_{MLE} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$$

---

<sup>1</sup>Bishop 2006.

- Assume now we have some prior  $p(\theta) = \mathcal{N}(\mu_\theta, \Sigma_\theta)$ .
- The posterior over  $\theta$  becomes

$$p(\theta^* | \mathbf{X}, \mathbf{y}) \propto p(\mathbf{y} | \mathbf{X}, \hat{f}_{\theta^*}) p(\theta^*) = \mathcal{N}(\hat{f}_\theta(\mathbf{X}), \sigma_y^2 \mathbf{I}) \mathcal{N}(\mu_\theta, \Sigma_\theta)$$

- Due to Gaussian conjugacy, we can solve this in closed form<sup>2</sup> as

$$p(\theta^* | \mathbf{X}, \mathbf{y}) = \mathcal{N}(\theta^*; \mu_\theta^*, \Sigma_\theta^*)$$

$$\text{with } \Sigma_\theta^* = \left( \Sigma_\theta^{-1} + \frac{1}{\sigma_y^2} \mathbf{X}^\top \mathbf{X} \right)^{-1}$$

$$\text{and } \mu_\theta^* = \Sigma_\theta^* \left( \Sigma_\theta^{-1} \mu_\theta + \frac{1}{\sigma_y^2} \mathbf{X}^\top \mathbf{y} \right)$$

---

<sup>2</sup>Bishop 2006.

- If we choose the prior to be zero-mean and isotropic, i.e.  
 $p(\theta) = \mathcal{N}(\mathbf{0}, \sigma_\theta^2 \mathbf{I})$ , the posterior simplifies to

$$p(\theta^* | \mathbf{X}, \mathbf{y}) = \mathcal{N}(\theta^*; \mu_\theta^*, \Sigma_\theta^*)$$

$$\text{with } \Sigma_\theta^* = \left( \frac{1}{\sigma_\theta^2} \mathbf{I} + \frac{1}{\sigma_y^2} \mathbf{X}^\top \mathbf{X} \right)^{-1}$$

$$\text{and } \mu_\theta^* = \frac{1}{\sigma_y^2} \Sigma_\theta^* \mathbf{X}^\top \mathbf{y}$$

- The posterior mode  $\mu_\theta^*$  is also called *maximum a posteriori (MAP) estimate*.
- Note that this is equivalent to ridge regression<sup>3</sup>

$$\mu_\theta^* = \arg \min_{\theta} \|\hat{f}_\theta(\mathbf{X}) - \mathbf{y}\|_2^2 + \lambda \|\theta\|_2^2 \quad \text{with } \lambda = \frac{\sigma_y^2}{\sigma_\theta^2}$$

---

<sup>3</sup>Bishop 2006.

- Assume we want to predict the response  $y^*$  at a new test point  $\mathbf{x}^*$ . The MAP estimate would just be  $\hat{y}^* = \mathbf{x}^{*\top} \boldsymbol{\mu}_\theta^*$ .
- In contrast, the full predictive posterior is

$$\begin{aligned} p(y^* | \mathbf{X}, \mathbf{y}, \mathbf{x}^*) &= \int p(y^* | \boldsymbol{\theta}^*, \mathbf{x}^*) p(\boldsymbol{\theta}^* | \mathbf{X}, \mathbf{y}) d\boldsymbol{\theta}^* \\ &= \int \mathcal{N}(\mathbf{x}^{*\top} \boldsymbol{\theta}^*, \sigma_y^2) \mathcal{N}(\boldsymbol{\mu}_\theta^*, \boldsymbol{\Sigma}_\theta^*) d\boldsymbol{\theta}^* \end{aligned}$$

- Due to Gaussian conjugacy this has a closed form<sup>4</sup> solution:

$$\begin{aligned} p(y^* | \mathbf{X}, \mathbf{y}, \mathbf{x}^*) &= \mathcal{N}(y^*; \boldsymbol{\mu}_{y^*}, \sigma_{y^*}^2) \\ \text{with } \boldsymbol{\mu}_{y^*} &= \mathbf{x}^{*\top} \boldsymbol{\mu}_\theta^* \\ \text{and } \sigma_{y^*}^2 &= \sigma_y^2 + \mathbf{x}^{*\top} \boldsymbol{\Sigma}_\theta^* \mathbf{x}^* \end{aligned}$$

---

<sup>4</sup>Murphy 2012.

- A Gaussian process is a prior over functions  $f \sim \mathcal{GP}(m(\cdot), k(\cdot, \cdot))$ , where  $m(\mathbf{x}) = \mathbb{E}[f(\mathbf{x})]$  is called a *mean function* and  $k(\mathbf{x}, \mathbf{x}') = \mathbb{E}[(f(\mathbf{x}) - m(\mathbf{x}))^\top (f(\mathbf{x}') - m(\mathbf{x}'))]$  is called a *kernel function*.
- The predictive posterior<sup>5</sup> is

$$p(\mathbf{y}^* \mid \mathbf{X}, \mathbf{y}, \mathbf{x}^*) = \mathcal{N}(\mu_y^*, \Sigma_y^*)$$

with  $\mu_y^* = \mathbf{k}(\mathbf{X}, \mathbf{x}^*)^\top [\mathbf{K}(\mathbf{X}, \mathbf{X}) + \sigma_y^2 \mathbf{I}]^{-1} \mathbf{y}$

and  $\Sigma_y^* = k(\mathbf{x}^*, \mathbf{x}^*) - \mathbf{k}(\mathbf{X}, \mathbf{x}^*)^\top [\mathbf{K}(\mathbf{X}, \mathbf{X}) + \sigma_y^2 \mathbf{I}]^{-1} \mathbf{k}(\mathbf{X}, \mathbf{x}^*) + \sigma_y^2 \mathbf{I}$

- Here,  $\mathbf{k}(\mathbf{X}, \mathbf{x}^*)$  is a column vector with elements  $\mathbf{k}(\mathbf{X}, \mathbf{x}^*)_i = k(\mathbf{x}_i, \mathbf{x}^*)$  and  $\mathbf{K}(\mathbf{X}, \mathbf{X})$  is a symmetric square matrix with elements  $\mathbf{K}(\mathbf{X}, \mathbf{X})_{ij} = k(\mathbf{x}_i, \mathbf{x}_j)$ .

---

<sup>5</sup>Rasmussen & Williams 2006.

- The kernel function  $k(\cdot, \cdot)$  has to be positive-semidefinite, i.e. the matrix  $\mathbf{K}(\mathbf{X}, \mathbf{X})$  has to be positive-semidefinite for every  $\mathbf{X}$ .
- Every positive-semidefinite kernel function can be written as an inner product in a certain feature (Hilbert) space<sup>6</sup>  $\mathcal{H}$ , i.e.  
$$k(\mathbf{x}, \mathbf{x}') = \Phi(\mathbf{x})^\top \Phi(\mathbf{x}')$$
 for some  $\Phi : \mathbb{R}^d \rightarrow \mathcal{H}$ .
- A popular kernel function is the *radial basis function* (RBF) kernel

$$k_{RBF}(\mathbf{x}, \mathbf{x}') = \exp\left(-\frac{1}{2\sigma^2}\|\mathbf{x} - \mathbf{x}'\|^2\right)$$

- It turns out that  $\dim(\mathcal{H}_{RBF}) = \infty$ . So the GP allows to perform a regression in an infinite-dimensional feature space, while the computational complexity only depends on  $n$ . It is thus a form of *nonparametric* regression.

---

<sup>6</sup>Mercer 1909.

- The Bayesian learning paradigm can be extended to complex parametric models, such as deep neural networks. This is called *Bayesian deep learning*.
- The basic assumptions are still the same, i.e.

$$p(\theta^* \mid \mathbf{X}, \mathbf{y}) \propto p(\mathbf{y} \mid \mathbf{X}, \hat{f}_{\theta^*}) p(\theta^*)$$
$$p(y^* \mid \mathbf{X}, \mathbf{y}, \mathbf{x}^*) = \int p(y^* \mid \theta^*, \mathbf{x}^*) p(\theta^* \mid \mathbf{X}, \mathbf{y}) d\theta^*$$

- However, without conjugacy, exact inference on the network weights  $\theta$  is usually not tractable anymore.
- One therefore needs to resort to approximate inference methods (variational inference, MCMC, etc.).

## Reinforcement learning

---

- There are generally two types of RL: *model-based* and *model-free*.
- Model-based RL involves learning a transition model  $P(s_{t+1} | s_t, a_t)$  of the environment, as well as a state value function  $V(s)$ .
- Model-free RL does not model the transitions, but directly learns a policy  $\pi(a_t | s_t)$  or a state-action value function (a.k.a. Q-function)  $Q(s_t, a_t)$ .
- Learning just the policy or Q-function can often be cheaper than learning the whole transition model.
- A good intuition for different types of value learning is provided at <https://distill.pub/2019/perspective-on-value-learning/>.

- In Q-learning, we learn a state-action value function  $Q(s_t, a_t)$ .
- The policy can then for instance be chosen to be  $\pi(a_t | s_t) = \delta(\arg \max_{a_t} Q(a_t, s_t))$ , where  $\delta(\cdot)$  is the Dirac measure.
- During learning, for every transition tuple  $(s_t, a_t, r_t, s_{t+1})$  we observe, we update the Q-function as

$$Q(s_t, a_t) \leftarrow (1 - \alpha_t)Q(s_t, a_t) + \alpha_t \left( r_t + \gamma \max_{a'} Q(s_{t+1}, a') \right)$$

- Here,  $\alpha_t$  is the step size and  $\gamma$  is the discount factor.

Questions?

## REFERENCES

- Bishop, C. M. (2006). Pattern recognition and machine learning. Springer Science+ Business Media.
- Murphy, K. P. (2012). Machine learning: a probabilistic perspective. MIT press.
- Rasmussen, C. E., & Williams, C. K. (2006). Gaussian Processes for Machine Learning. MIT press.
- Mercer, J. (1909). Functions of positive and negative type, and their connection with the theory of integral equations. Proceedings of the Royal Society of London. Series A, Containing Papers of a Mathematical and Physical Character, 83(559), 69-70.