

TUTORIAL

PROBABILISTIC ARTIFICIAL INTELLIGENCE

REINFORCEMENT LEARNING

CHARLOTTE BUNNE

ETH ZURICH

WINTER TERM 2019

Overview:

Part 0: MDPs and Bellman Equations

Part 1: Model-Based Reinforcement Learning

Part 2: Model-Free Reinforcement Learning

Part 2: (Cool) Example of Q-Learning

Part 3: Discussion of Homework


MDPS AND BELLMAN EQUATIONS

MARKOV DECISION PROCESS (MDPs)

A MDP is defined by ...

- ... a set of states $X = \{1 \dots n\}$
- ... a set of actions $A = \{1 \dots m\}$
- ... transition probabilities $p(x'|x, a)$
- ... a reward function r

Planning via a (deterministic) policy
 $\pi : X \rightarrow A$



Several optimality criteria exist:

Finite Horizon: $\mathbb{E} \left(\sum_{t=0}^T r_t \right)$

Average Reward: $\lim_{h \rightarrow \infty} \mathbb{E} \left(\frac{1}{h} \sum_{t=0}^{\infty} r_t \right)$

Infinite Horizon: $\mathbb{E} \left(\sum_{t=0}^{\infty} \gamma^t r_t \right)$

- ▷ discount factor $0 \leq \gamma \leq 1$ which quantifies importance of future rewards

Find the optimal policy via $V^*(x) = \max_{\pi} \mathbb{E} (\sum_{t=0}^{\infty} \gamma^t r_t)$ (optimality criterium)

Bellman equation (1957) defines *recursively* the **value function** of policy π at state x :

$$V^{\pi}(x) = r(x, \pi(x)) + \gamma \sum_{x'} p(x'|x, \pi(x)) V^{\pi}(x')$$

▷ Dynamic Programming

Assumption: transition probabilities and reward function are known

BELLMAN THEOREM

Value Function V^π

$$V^\pi(x) = r(x, \pi(x)) + \gamma \sum_{x'} p(x'|x, \pi(x)) V^\pi(x')$$

↑

Greedy Policy w.r.t. V

$$\arg \max_a r(x, a) + \gamma \sum_{x'} p(x'|x, a) V(x')$$

VALUE ITERATION ALGORITHM

Value iteration computes the optimal state value function by iteratively improving estimate of $V(x)$

Algorithm 1 Pseudocode for Value Iteration Algorithm


```
1: input  $A, X, r, p$ 
2: initialize  $V_0$  to arbitrary values
3: while not converged do
4:   for  $x \in X$  do
5:     for  $a \in A$  do
6:        $Q_t(x, a) \leftarrow r(x, a) + \gamma \sum_{x'} p(x'|x, a)V_{t-1}(x')$ 
7:     end for
8:      $V_t(x) \leftarrow \max_a Q_t(x, a)$ 
9:   end for
10: end while
11: return  $V_T$ 
```

REINFORCEMENT LEARNING

= Planning in unknown MDPs

Challenges:

- Exploration-Exploitation Dilemma
- Credit Assignment Problem
- Function Approximation



DQN [Minh et al. (2015)]

REINFORCEMENT LEARNING

Approaches:

Model-Based Reinforcement Learning

Model-Free Reinforcement Learning

On-Policy Reinforcement Learning

Off-Policy Reinforcement Learning

EXPLORATION-EXPLOITATION DILEMMA

Common Approaches:

1. ϵ -Greedy

$$\pi = \begin{cases} \arg \max_a r(a), & \text{with prob. } 1 - \epsilon \\ \text{rand}(a), & \text{with prob. } \epsilon \end{cases}$$

EXPLORATION-EXPLOITATION DILEMMA

Common Approaches:

1. ϵ -Greedy

$$\pi = \begin{cases} \arg \max_a r(a), & \text{with prob. } 1 - \epsilon \\ \text{rand}(a), & \text{with prob. } \epsilon \end{cases}$$

- ▷ if ϵ satisfies **Robbins Monro (RM)** conditions ($\sum_{t=1}^{\infty} \epsilon_t = \infty$ and $\sum_{t=1}^{\infty} \epsilon_t^2 < \infty$), it converges to optimal policy with probability 1

EXPLORATION-EXPLOITATION DILEMMA

Common Approaches:

1. ϵ -Greedy

$$\pi = \begin{cases} \arg \max_a r(a), & \text{with prob. } 1 - \epsilon \\ \text{rand}(a), & \text{with prob. } \epsilon \end{cases}$$

- ▷ if ϵ satisfies **Robbins Monro (RM)** conditions ($\sum_{t=1}^{\infty} \epsilon_t = \infty$ and $\sum_{t=1}^{\infty} \epsilon_t^2 < \infty$), it converges to optimal policy with probability 1
- ▷ performs quite well **but** does not quickly eliminate clearly suboptimal actions

EXPLORATION-EXPLOITATION DILEMMA

Common Approaches:

1. ϵ -Greedy

$$\pi = \begin{cases} \arg \max_a r(a), & \text{with prob. } 1 - \epsilon \\ \text{rand}(a), & \text{with prob. } \epsilon \end{cases}$$

2. Optimistic Initialization: *optimism under uncertainty*

EXPLORATION-EXPLOITATION DILEMMA

Common Approaches:

1. ϵ -Greedy

$$\pi = \begin{cases} \arg \max_a r(a), & \text{with prob. } 1 - \epsilon \\ \text{rand}(a), & \text{with prob. } \epsilon \end{cases}$$

2. Optimistic Initialization: *optimism under uncertainty*
3. Upper Confidence Bound (UCB)
4. Posterior Sampling

MODEL-BASED REINFORCEMENT LEARNING

MODEL-BASED REINFORCEMENT LEARNING

Strategy: Learn MDP by estimating the transition probabilities and reward function from observed variables $\langle x, a, r, x' \rangle$

R-MAX ALGORITHM

1. Start with an optimistic MDP of the environment
 - ▷ *optimism under uncertainty*: initialize $r = R_{max}$ and $p(x^*|x, a) = 1$ for $\forall x \in X, a \in A$, with x^* being a fictitious state

R-MAX ALGORITHM

1. Start with an optimistic MDP of the environment

▷ *optimism under uncertainty*: initialize $r = R_{max}$ and $p(x^*|x, a) = 1$ for $\forall x \in X, a \in A$,
with x^* being a fictitious state
implicit explore or exploit approach

R-MAX ALGORITHM

1. Start with an optimistic MDP of the environment
 - ▷ *optimism under uncertainty*: initialize $r = R_{max}$ and $p(x^*|x, a) = 1$ for $\forall x \in X, a \in A$, with x^* being a fictitious state
2. Solve for optimal policy π in optimistic MDP (for given r and p)

R-MAX ALGORITHM

1. Start with an optimistic MDP of the environment
 - ▷ *optimism under uncertainty*: initialize $r = R_{max}$ and $p(x^*|x, a) = 1$ for $\forall x \in X, a \in A$, with x^* being a fictitious state
2. Solve for optimal policy π in optimistic MDP (for given r and p)
3. Take greedy action according to policy π
4. Update MDP by estimating p and updating r
5. Repeat (from 2)

R-MAX ALGORITHM

1. Start with an optimistic MDP of the environment
 - ▷ *optimism under uncertainty*: initialize $r = R_{max}$ and $p(x^*|x, a) = 1$ for $\forall x \in X, a \in A$, with x^* being a fictitious state
2. Solve for optimal policy π in optimistic MDP (for given r and p)
3. Take greedy action according to policy π
4. Update MDP by estimating p and updating r
5. Repeat (from 2)

Implicit explore or exploit approach

Agent acts greedily according to a model that assumes all *underexplored* states are maximally rewarding

Definition: Probably Approximately Correct (PAC)-MDP¹

An algorithm \mathcal{A} is said to be an efficient PAC-MDP algorithm if, for any $\epsilon > 0$ ('accuracy'), $0 < \delta < 1$ ('probability of failure'), the per-timestep computational complexity, space complexity and sample complexity of \mathcal{A} are less than some polynomial in the relevant quantities $(X, A, \frac{1}{\delta}, \frac{1}{(1-\gamma)})$, with probability at least $1 - \delta$.

Result: R-max is a PAC-MDP algorithm, it is called *efficient* because its sample complexity scales only polynomially in the number of states

¹[Strehl, Li, Littman (2009)]

Definition: Probably Approximately Correct (PAC)-MDP¹

An algorithm \mathcal{A} is said to be an efficient PAC-MDP algorithm if, for any $\epsilon > 0$ ('accuracy'), $0 < \delta < 1$ ('probability of failure'), the per-timestep computational complexity, space complexity and sample complexity of \mathcal{A} are less than some polynomial in the relevant quantities $(X, A, \frac{1}{\delta}, \frac{1}{(1-\gamma)})$, with probability at least $1 - \delta$.

Result: R-max is a PAC-MDP algorithm, it is called *efficient* because its sample complexity scales only polynomially in the number of states

- ▷ In natural environments, however, number of states is enormous: it is exponential in the number of objects in the environment

¹[Strehl, Li, Littman (2009)]

MODEL-FREE REINFORCEMENT LEARNING

MODEL-FREE REINFORCEMENT LEARNING

How can we solve Bellman's fixed point equation without knowing transition probabilities and reward?

$$V^*(x) = \max_a r(x, a) + \gamma \sum_{x'} p(x'|x, a) V^*(x')$$

$$\pi^*(x) \in \arg \max_a r(x, a) + \underbrace{\gamma \sum_{x'} p(x'|x, a) V^*(x')}_{Q^*(x, a)}$$

Estimate V or Q directly from samples (via stochastic approximation methods).

TEMPORAL DIFFERENCE (TD) LEARNING

Observe samples $\langle x, a, r, x' \rangle$

Assumption: if value estimates are correct, the following must hold

$$V(x) = r + \gamma V(x')$$

Otherwise there is an error: $\delta = r + \gamma V(x') - V(x)$ (TD error)

To learn a better estimate minimize δ and with learning rate α update

$$V(x) \leftarrow V(x) + \alpha \underbrace{(r + \gamma V(x') - V(x))}_{\text{TD target}}$$

Q-LEARNING

Q-function ($Q : X \times A \rightarrow \mathbb{R}$) quantifies quality of a state-action combination

Core of the algorithm: **value-iteration update**

- ▷ weighted average of old value and new information

$$\begin{aligned} Q(x, a) &= Q(x, a) + \alpha \left(r + \gamma \max_{a'} Q(x', a') - Q(x, a) \right) \\ &= (1 - \alpha) \underbrace{Q(x, a)}_{\text{old value}} + \alpha \overbrace{\left(r + \gamma \underbrace{\max_{a'} Q(x', a')}_{\text{estimate of optimal future value}} \right)}^{\text{learned value}} \end{aligned}$$

Stochastic Approximation used in Q-learning and TD-learning to estimate the value function

- ▷ Q-Learning: (asynchronous) implementation of the Robbins Monro algorithm for finding fixed points
- ▷ To prove convergence of Q-Learning we need results from Robbins Monro

$$\sum_{t=1}^{\infty} \alpha_t = \infty \quad \text{and} \quad \sum_{t=1}^{\infty} \alpha_t^2 < \infty$$

ROBBINS MONRO ALGORITHM

A classical methodology, studied by Robbins and Monro is

$$\theta_n = \theta_{n-1} - \alpha_n (h(\theta_{n-1}) + \varepsilon_n)$$

where $h : \mathbb{R}^d \rightarrow \mathbb{R}^d$, $(\varepsilon_n)_{n \geq 1}$ are i.i.d. random variables²

- ▷ Stochastic Gradient Descent (SGD)

Robbins-Munro: general sufficient conditions for iterates to converge,

$$\sum_{t=1}^{\infty} \alpha_t = \infty : \text{ convergence can be towards any point in } \mathbb{R}^d$$

$$\sum_{t=1}^{\infty} \alpha_t^2 < \infty : \text{ guarantees convergence}$$

²[Robbins & Munro (1951)]

Algorithm 2 Pseudocode for Q-Learning Algorithm

```
1: input actions, states, rewards, learning rate  $\alpha$ , discount factor  $\gamma$ 
2: initialize  $Q : X \times A \rightarrow \mathbb{R}$ 
3: for each episode do
4:   observe initial state  $x$ 
5:   for each step  $t = 0, 1, 2, \dots$  do
6:     take action  $a$  and observe  $r$  and  $x'$ 
7:      $Q(x, a) \leftarrow (1 - \alpha)Q(x, a) + \alpha(r + \gamma \max_{a'} Q(x', a'))$ 
8:      $x = x'$ 
9:   end for
10: end for
11: return  $Q$ 
```

[Watkins (1989), Dayan & Watkins (1992)]

Q-Learning with Function Approximation

To generalize over states and actions, parametrize Q with a function approximation, e.g. deep neural networks

$$\delta = r + \gamma \max_{a'} Q(x', a', \theta) - Q(x, a, \theta)$$

Learning is unstable due to **deadly triad** [Sutton & Barto (2018)]

- ▷ off-policy learning, flexible function approximation, bootstrapping

First stable Q-learning with function approximation: DQN [Minh et al. (2015)]

SOLUTION OF HOMEWORK

REFERENCE

Slides based on material accompanying the textbook 'AI: A Modern Approach' (3rd edition) by S. Russell and P. Norvig and the **NeuRIPS 2019 Tutorial** by Katja Hofmann.